
ASO Foundry and Sales AI Agents Standard

Sales AI agent responsibility, Foundry orchestration, code structure, canonical contracts, safety policy, and evaluation approach.

GitHub path: `foundry/ASO_Foundry_and_Sales_AI_Agents_Standard.md`

Version 1.1 - public customer-ready split package - refined for readability and GitHub publishing

Safety note: this public reference pack must not contain tenant IDs, credentials, secrets, real customer data, private endpoints, or connection strings.

Foundry and Dynamics 365 Sales AI agent standard

This file covers two AI layers: Dynamics 365 Sales AI/Copilot Sales agents for seller-facing intelligence, and Microsoft Foundry for parent orchestration, child agents, safety, routing, schema validation, and deterministic handoff back to Power Automate.

Sales AI / Copilot Sales agent responsibility

Agent	Primary responsibility	Must not do
Sales Qualification Agent	Research and assess inbound leads, fit, routing support, and seller-facing rationale.	No ERP posting; no direct customer lifecycle sends; no bypass of Foundry policy.
Sales Opportunity / Deal Close Agent	Assess opportunity importance, risk, close-readiness, and recommendations.	No SAP posting; no workflow execution; no direct lifecycle sends.

Sales agent configuration standards

Configuration area	Standard
Mode	Start with Research-only or restricted engagement. Research and engage requires explicit compliance and Customer Insights boundary review.
Knowledge sources	Use approved public sites and approved internal documents only. Do not add unmanaged SAP exports, confidential pricing, contract templates, or uncontrolled sensitive files.
Lead selection	Pilot only: Open leads, approved source, email exists, consent not blocked, not already Running, pilot owner/BU.
Opportunity selection	Pilot only: Open opportunities, active stages, close date in pilot window, excludes closed won/lost and Running records.
Writeback	Native agent output may be seller-visible only. Foundry/Power Automate normalizes governed fields where native writeback is limited.
Testing	Use 5-10 pilot leads/opportunities. Confirm no unauthorized customer message and no ERP action occurred.

Foundry responsibilities

Foundry owns	Foundry does not own
End-to-end orchestration visibility; agent sequencing and routing; safety guardrails; SAP-aware tool coordination through APIM; schema validation; separation of CRM facts, SAP facts, sales-agent outputs, and inferences; deterministic handoff to Power Automate.	CRM transactional truth; ERP posting; customer communication sends; secret storage; seller decision rights; unreviewed commercial actions.

Repository structure standard

```
aso-foundry-orchestrator/  
  src/  
    orchestrator/  
      parent.py  
      router.py  
      policy.py  
      schemas.py  
      telemetry.py
```

```

agents/
  lead_origination.py
  nurturing.py
  qualification.py
  opportunity_classification.py
  risk_competitor.py
  next_best_action.py
  handoff.py
  account_development.py
adapters/
  dataverse_client.py
  sap_client.py
  sales_agent_context.py
  customer_insights_context.py
contracts/
  canonical_input.schema.json
  canonical_output.schema.json
  child_outputs.schema.json
tests/
  unit/
  integration/
  evaluation/
pipelines/
README.md

```

Canonical input and output requirements

Contract element	Required fields
Input identity	message_id, correlation_id, record_type, record_id, trigger_type, trigger_reason
Tenant context	environment_name, business_unit, initiating_system
External refs	hubspot_contact_id, sap_customer_id, sap_business_partner_id, customer_insights_profile_id
Sales agent context	Qualification status, score, rationale; opportunity status, risk, recommendation
Execution policy	max_agent_retries, allow_sap_reads, allow_sap_writes, allow_customer_insights_activation, human_approval_required
Output identity	message_id, correlation_id, status, summary, record_type, record_id, trace_id
Output actions	update_dataverse, create_task, notify, submit_for_approval, enqueue_retry, start_customer_journey, sync_consent_state, none
Output governance	confidence, requires_human_review, recommended_lifecycle_stage, facts grouped by CRM/SAP/sales_agent/inferred

Safety policy

- Block or escalate when confidence is below threshold.
- Escalate when Sales Qualification Agent and Foundry disagree materially.
- Escalate when SAP account match is ambiguous.

-
- Escalate when routing changes ownership or account ownership.
 - Block communication when it violates consent state.
 - Escalate high-risk opportunity commercial actions.
 - Block ERP submit if approval is missing.
 - Reject output JSON that fails schema validation.